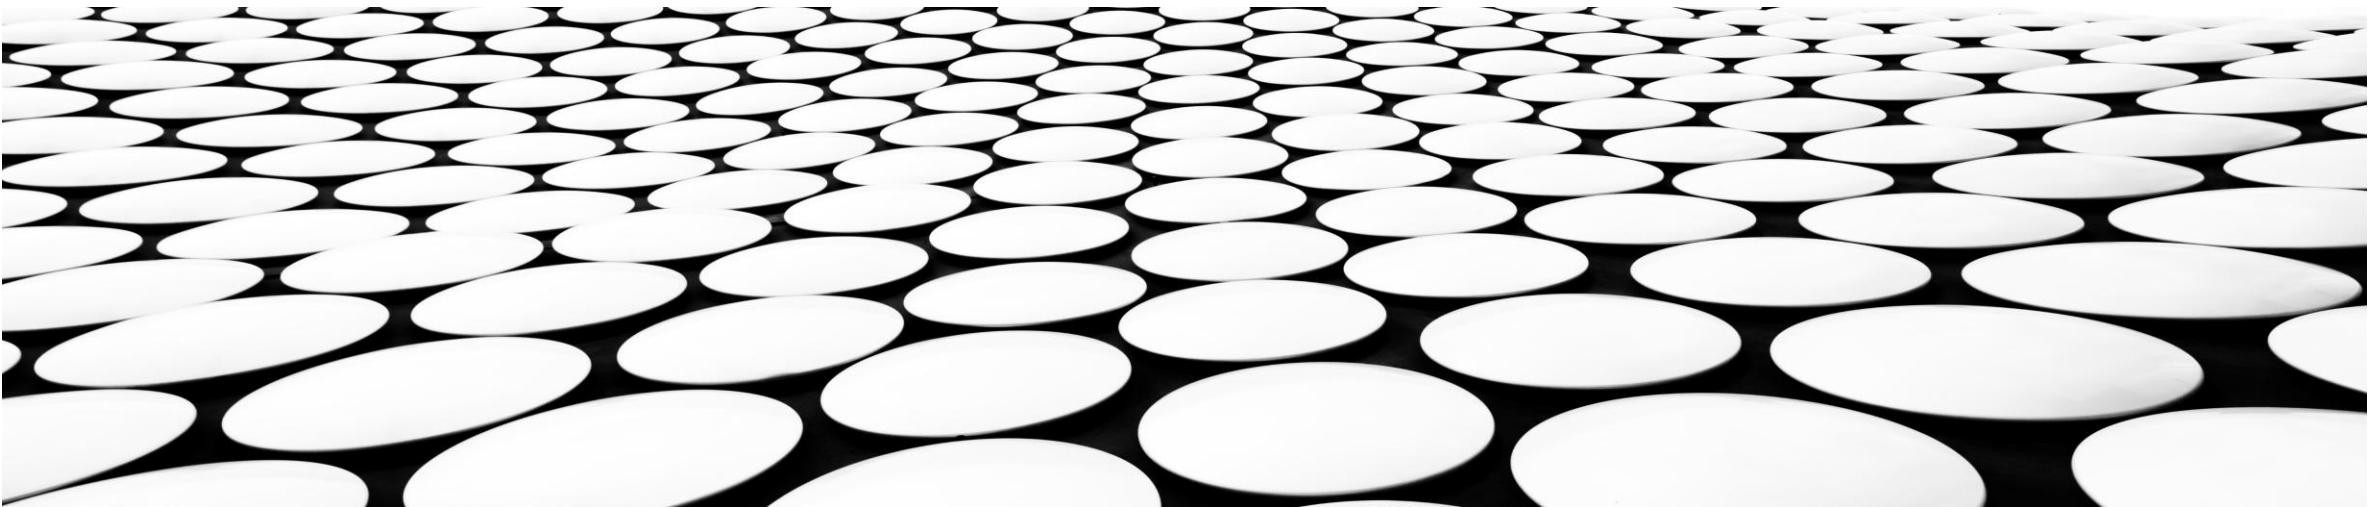

DATA MINING AND PREDICTIVE DATA ANALYSIS

CHAPTER-2

DATA PREPROCESSING



INTRODUCTION TO DATA PREPROCESSING

- **Data in the real world are - Noisy, Incomplete and Inconsistent**
 - Due to their typically huge size (often several gigabytes or more).
 - As their likely origin from multiple, heterogeneous sources
 - Historical/past un-attended data- often contains field values that have expired, are no longer relevant, or are simply missing.
- Real-world (raw data) data is often:
 - **Incomplete:** missing values, lacking attributes, or containing only partial information.
 - **Noisy:** containing errors, outliers, or irrelevant information.
 - **Inconsistent:** conflicting values, different units/formats, or duplication.
 - **High-dimensional:** too many features, making analysis complex.

WHY DATA PREPROCESSING

- **No quality data, no quality mining results!**

- Quality decisions must be based on quality data.
- “garbage in, garbage out”: if raw data is noisy, incomplete, or inconsistent, the mining results will be poor.

- **Data Preprocessing**

- It prepares raw data into a clean, consistent, and suitable format for further analysis.
- Is the first and most critical step in the data mining process.
- Improve the quality of the data and, consequently the quality of the mining results and lead to huge payoffs for decision making.
- Data preprocessing can take 20–60% of the total time and effort in the data mining process based on the size and complexity of the data set.

WHY DATA PREPROCESSING

- **Importance of Data Preprocessing**
 - Ensures better quality results (reliable, efficient, and meaningful mining results)
 - Reduces computational cost and time.
 - Increases accuracy of models (classification, clustering, prediction).
- In order to be useful for data mining purposes, the databases need to undergo preprocessing, in the form of
 - **Data Cleaning**, and
 - **Data Transformation**

DATA CLEANING

- let us take a look at some of the kinds of errors that could creep into even a tiny data set

Customer ID	Zip	Gender	Income	Age	Marital Status	Transaction Amount
1001	10048	M	75,000	C	M	5000
1002	J2S7K7	F	-40,000	40	W	4000
1003	90210		10,000,000	45	S	7000
1004	6269	M	50,000	0	S	1000
1005	55101	F	99,999	30	D	3000

- Without proper cleaning, data mining algorithms may produce misleading patterns- leading to inaccurate, inconsistent, and unreliable outcomes

DATA CLEANING

- Data cleaning (also called **data cleansing** or **data scrubbing**) is the process of detecting and correcting (or removing) corrupt, inaccurate, incomplete, or irrelevant records from a dataset.
- Causes of Dirty Data
 - Human errors – Typos, inconsistent spellings, missing entries.
 - Data integration issues – Different formats, units, or coding standards across sources.
 - Noise – Random variations or outliers in data.
 - Missing values – Due to non-response, equipment failure, or improper data entry.
 - Redundancy and duplicates – Same data recorded multiple times.

DATA CLEANING

■ Types of Data Quality Issues

- Missing Data: Null values, blanks, unknowns.
- Noisy Data: Errors, outliers, misrecorded values.
- Inconsistent Data: Conflicting values (e.g., "Male" vs. "M").
- Duplicate Data: Repeated records.
- Irrelevant Data: Attributes not useful for analysis.

HANDLING MISSING DATA

- **Missing Values:** Some tuples may not have recorded value for several attributes.
- A common method of “handling” missing values is simply to omit the records or fields with missing values -- However, this may be dangerous
 - Deleting the records with missing values would lead to a biased subset of the data.
 - Wasteful to omit the information in all other fields, just because one field value is missing.
- This method may be effective, if the tuple contains several attributes with missing values.
- This is usually done when the class label is missing in classification problems.

HANDLING MISSING DATA

☐ Replace the missing value with some constant, specified by the analyst.

- Missing numeric values replaced with 0.0
- Missing categorical values replaced with “Missing”

	mpg	cylinders	cubicinches	hp
1	14.00		350.00	165.00
2	31.90		89.00	71.00
3	517.00		302.00	140.00
4	15.00	Missing	400.00	150.00
5	30.50	Missing	0.00	0.00
6	23.00	Missing	350.00	125.00
7	13.00	Missing	351.00	158.00
8	14.00		0.00	215.00
9	25.40		0.00	77.00
10	37.70		89.00	62.00

HANDLING MISSING DATA

□ Replace Missing Values with Mode or Mean

- Attribute mean- for numeric variables
- Attribute mode -for categorical variables

- Mode of categorical field cylinders = 4
- Mean for non-missing values in numeric field cubicinches = 200.65

	mpg	cylinders	cubicinches	hp
	continuous	categorical	continuous	continuous
1	14.00		350.00	165.00
2	31.90		89.00	71.00
3	517.00		302.00	140.00
4	15.00		400.00	150.00
5	30.50		200.65	106.53
6	23.00		350.00	125.00
7	13.00		351.00	158.00
8	14.00		200.65	215.00
9	25.40		200.65	77.00
10	37.70		89.00	62.00

HANDLING MISSING DATA

- Substituting mode or mean for missing values sometimes works well.
- Mean not always best choice for “typical” value.
- Resulting confidence levels for statistical inference become overoptimistic.
- Domain experts should be consulted regarding approach to replace missing values.
- Benefits and drawbacks resulting from the replacement of missing values must be carefully evaluated.

HANDLING MISSING DATA

Replace Missing Values with Random

- Replace the missing values with a value generated at random from the observed distribution of the variable.
 - Value for cylinders, cubicinches, and hp randomly drawn proportionately from each field's distribution
 - Method superior compared to mean substitution
 - Measures of location and spread remain closer to original

Table View (17)					
	mpg	cylinders	cubicinches	hp	
	continuous	categorical	continuous	continuous	
1	14.00	8	350.00	165.00	
2	31.90	4	89.00	71.00	
3	517.00	8	302.00	140.00	
4	15.00	8	400.00	150.00	
5	30.50	4	144.15	116.55	
6	23.00	4	350.00	125.00	
7	13.00	6	351.00	158.00	
8	14.00	8	323.45	215.00	
9	25.40	5	81.84	77.00	
10	37.70	4	89.00	62.00	

No guarantee resulting records make sense

HANDLING MISSING DATA

❑ Replace Missing Values using - Data Imputation Method

- Replace the missing values with imputed values based on the other characteristics of the record.
- What would be the most likely value for this missing value - given all the other attributes for a particular record
- Example:
 - Regression-based imputation.
 - Using k-Nearest Neighbors (kNN) for estimation.
 - inference-based tools using a Bayesian inference, or decision tree

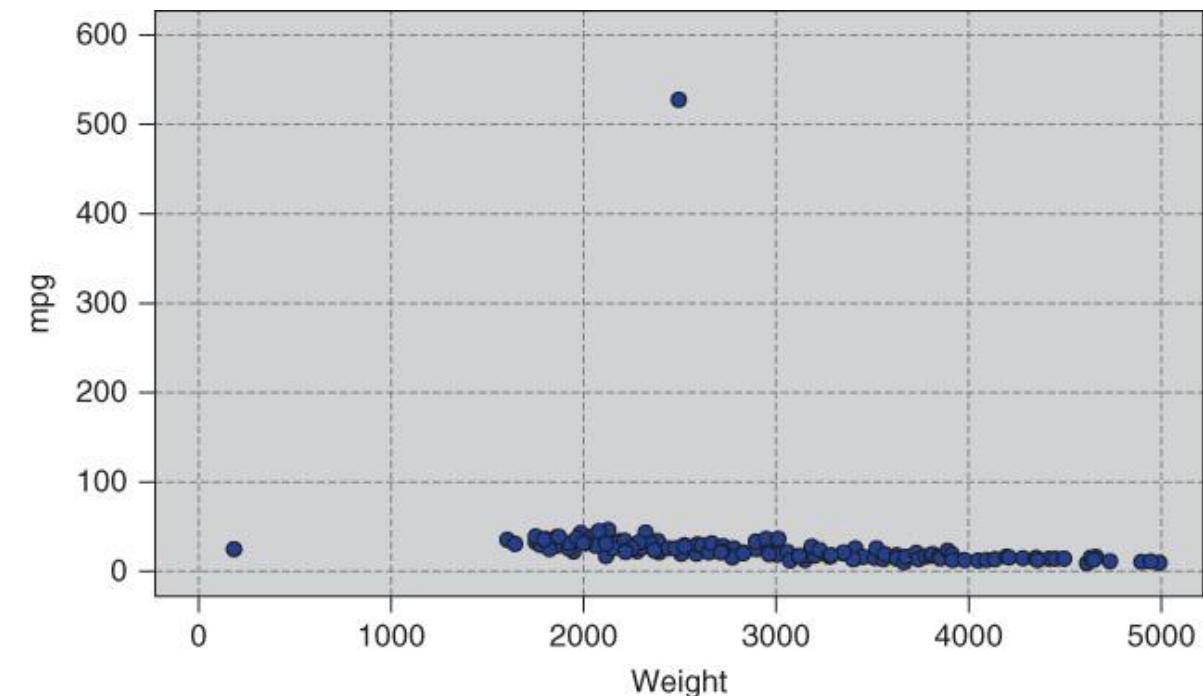
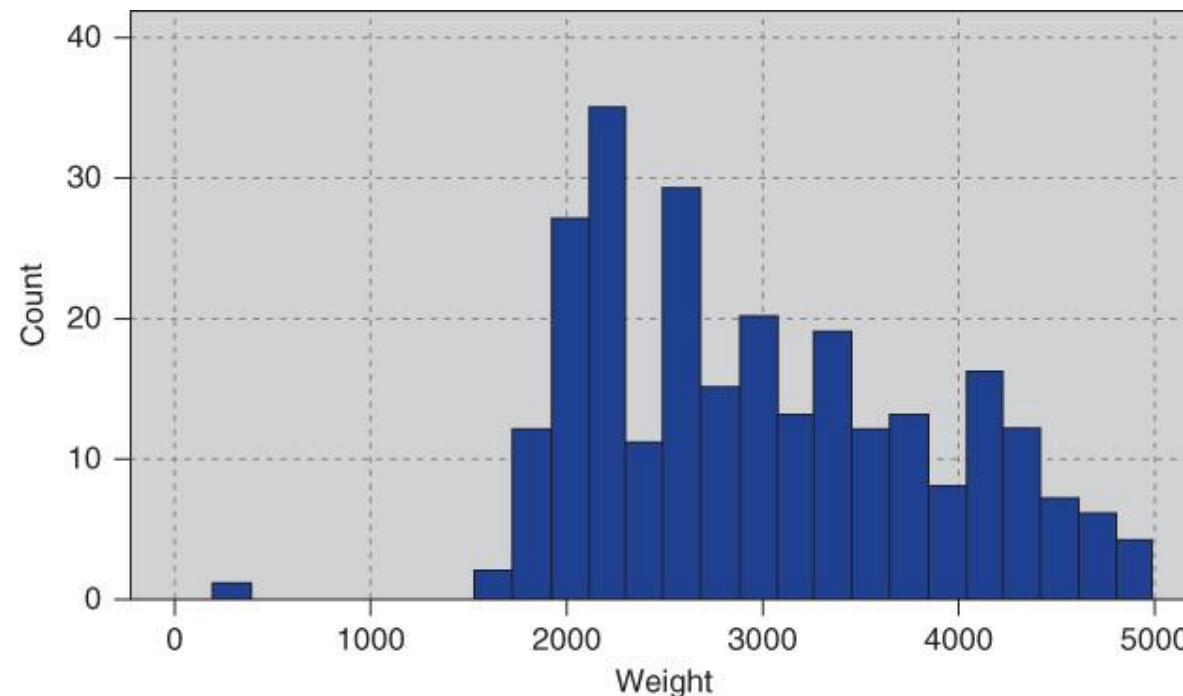
GRAPHICAL METHODS FOR IDENTIFYING OUTLIERS

- **Outlier :** A data point that significantly deviates from the overall pattern of the dataset.
- Outliers are extreme values that go against the trend of the remaining data.
- Outliers may arise due to:
 - Errors in data collection or entry.
 - Natural variation in data.
 - Rare events or anomalies worth studying (Outlier Analysis)
- Identifying outliers is important because they may that can distort analysis..
- Statistical methods are sensitive to the presence of outliers

GRAPHICAL METHODS FOR IDENTIFYING OUTLIERS

- Graphical Method For Identifying Outliers

- Histogram - for numeric variables
- Scatter Plot - to reveal outliers in more than one variable



MEASURES OF CENTER AND SPREAD

- **Central Tendency:** Central tendency refers to a single value that attempts to describe a dataset by identifying the central point within that dataset.
- Purpose in Data Mining:
 - Provides quick insight into the distribution of values.
 - Helps detect skewness and outliers.
 - Serves as a baseline for further statistical analysis and preprocessing.
- Most fundamental descriptive measures for “central tendency” of data-
Mean, Median, and Mode,

MEASURES OF CENTER AND SPREAD

- **Mean (Arithmetic Average):** The sum of all data values divided by the number of values.
 - Example: Dataset = {2, 4, 6, 8, 10} Mean = $(2 + 4 + 6 + 8 + 10)/5 = 30/5 = 6$
 - Advantages: Easy to compute, uses all data points.
 - Limitations: Sensitive to outliers (e.g., extreme values can distort the mean).
- **Median:** The middle value when the data is arranged in ascending (or descending) order.
 - Calculation: If n is odd \rightarrow Median = middle value.
 - If n is even \rightarrow Median = average of the two middle values.
 - Example: Dataset = {3, 5, 7, 9, 11} \rightarrow Median = 7
Dataset = {3, 5, 7, 9} \rightarrow Median = $(5+7)/2 = 6$
 - Advantages: Not affected by outliers or skewed data.
 - Limitations: Ignores the magnitude of values (uses only positional data).

MEASURES OF CENTER AND SPREAD

- **Mode:** The value(s) that appear most frequently in a dataset.
 - Types:
 - Unimodal: One mode
 - Bimodal: Two modes
 - Multimodal: More than two modes
 - No mode: If all values occur with the same frequency
 - Example: Dataset = {2, 4, 4, 6, 8, 8, 8, 10} → Mode = 8
 - Advantages: Useful for categorical data (e.g., "most purchased product").
 - Limitations: May not exist or may not be unique.

MEASURES OF CENTER AND SPREAD

Measure	Best Used When	Sensitive to Outliers?	Data Type
Mean	Symmetric, numeric data	Yes	Interval/Ratio
Median	Skewed data or with outliers	No	Ordinal, Interval/Ratio
Mode	Most frequent item needed	No	Nominal, Ordinal, Interval

DATA TRANSFORMATION

■ Data Transformation

- Refers to a set techniques used to transform raw data into a more appropriate format to enable efficient data mining and model building.
- In data transformation, the raw data are transformed or consolidated into forms appropriate for mining.
- With data transformation,
 - The resulting mining process becomes more efficient and faster
 - Enhances compatibility in model building
 - The patterns found are of better quality and may be easier to understand.

DATA TRANSFORMATION

- **Strategies for data transformation include the following:**

1. **Data Smoothing:** works to remove noise from the data. Techniques include binning, regression, and clustering.
2. **Attribute construction:** new attributes are constructed and added from the given set of attributes to help the mining process.
3. **Aggregation:** summary or aggregation operations applied to the data.
 - For example: The *daily sales* data - aggregated to - *monthly* and *annual sales* amounts.
 - Typically used in constructing a **data cube** for data analysis

DATA TRANSFORMATION

4. Normalization: the attribute data are scaled so as to fall within a smaller range, such as -1.0 to 1.0, or 0.0 to 1.0.

5. Discretization: the raw values of a numeric attribute (e.g., age) are replaced by interval labels (e.g., 0–10, 11–20, etc.) or conceptual labels (e.g., *youth*, *adult*, *senior*).

- The labels, in turn, can be recursively organized into higher-level concepts, resulting in a *concept hierarchy* for the numeric attribute.

DATA TRANSFORMATION BY NORMALIZATION

Data Normalization

- The measurement unit used can affect the data analysis.
 - Attributes tend to have ranges different from each other
 - Expressing an attribute in smaller units will lead to a larger range, and thus tend to give such an attribute greater effect or “weight.”
- Variables with greater ranges tend to have larger influence on data model’s results.
- To help avoid dependence on the choice of measurement units (and thereby - the range), the data (numeric data) should be *normalized* or *standardized*.

DATA TRANSFORMATION BY NORMALIZATION

- Normalizing the data attempts to give all attributes an equal weight.
 - Particularly useful for classification algorithms involving *neural networks* or distance measurements such as *nearest-neighbor classification* and *clustering*.
- For distance-based methods, normalization helps prevent attributes with initially large ranges (e.g., *income*) from outweighing attributes with initially smaller ranges (e.g., binary attributes).

DATA TRANSFORMATION BY NORMALIZATION

❑ Methods for data normalization.

- There are several methods for data normalization.
 - *min-max normalization,*
 - *z-score normalization,*
 - *normalization by decimal scaling.*

DATA TRANSFORMATION BY NORMALIZATION

□ Min-max normalization

- It performs a linear transformation on the original data.
- Min-max normalization maps a value, v , of A to v' in the range $[new_min_A, new_max_A]$ by computing:

$$v' = \frac{v - min_A}{max_A - min_A} (new_max_A - new_min_A) + new_min_A$$

- min_A : minimum values of an attribute, A .
- max_A : maximum values of an attribute, A .

DATA TRANSFORMATION BY NORMALIZATION

- Example (Min-max normalization)

- Suppose that the minimum and maximum values for the attribute *income* are \$12,000 and \$98,000, respectively.
- Map income to the range [0.0, 1.0].
- By min-max normalization, a value of \$73,600 for income is transformed to

$$\frac{73,600 - 12,000}{98,000 - 12,000} (1.0 - 0) + 0 = 0.716$$

- Min-max normalization preserves the relationships among the original data values. It will encounter an “out-of-bounds” error if a future input case for normalization falls outside of the original data range for A

DATA TRANSFORMATION BY NORMALIZATION

□ z-score normalization

- The values for an attribute are normalized based on the *mean* and *standard deviation*.
- A value, v , for an attribute A , is normalized to v' by computing:

$$v' = \frac{v - \bar{A}}{\sigma_A}$$

- \bar{A} = mean of A
- σ_A = standard deviation of A
- This method of normalization is useful when the *actual minimum* and *maximum* of attribute A are unknown.

DATA TRANSFORMATION BY NORMALIZATION

- Example (z-score normalization)
 - Suppose for the attribute *income*
mean = \$54,000
standard deviation = \$16,000

A value of \$73,600 for *income* is transformed to

$$\frac{73,600 - 54,000}{16,000} = 1.225$$

Note:

- Data values that lie below the mean will have a negative Z-score standardization.
- Data values that lie above the mean will have a positive Z-score standardization
- Data values falling exactly on the mean will have a Z-score standardization of zero

DATA TRANSFORMATION BY NORMALIZATION

□ Normalization by Decimal Scaling

- Normalizes by moving the decimal point of values of attribute A .
- The number of decimal points moved depends on the *maximum absolute value* of A .
- A value, v , of A is normalized to v' by computing

$$v' = \frac{v}{10^j}$$

- Where, j is the smallest integer such that $\text{Max}(|v'|) < 1$.
- Decimal scaling ensures that every normalized value lies between -1 and 1

DATA TRANSFORMATION BY NORMALIZATION

- Example (Normalization by Decimal Scaling)
 - Suppose that the recorded values of A range from -986 to 917.
 - The maximum absolute value of A is 986.
 - To normalize by decimal scaling, we therefore divide each value by 1,000 (i.e., $j = 3$) so that -986 normalizes to -0.986 and 917 normalizes to 0.917.

FLAG VARIABLES

■ Flag Variables (Dummy Variables Or Indicator Variables)

- In many analytical methods (e.g., regression, classification), predictors must be numeric.
- However, real-world data often contains categorical variables.
- To use them in models, we convert categories into numeric form using flag variables.
- **Definition:** A **flag variable** is a binary variable that takes values:
 - 1 → if the condition is satisfied (category present)
 - 0 → if the condition is not satisfied (category absent)
- **Example:** Smoker_Flag
 - Smoker = Yes → 1 , Smoker = No → 0

FLAG VARIABLES

■ Handling Multiple Categories with Flags

- If a categorical variable has **k categories**:
 - We typically create **$k - 1$ flag variables**.
 - The unassigned category becomes the **reference category**.
 - This avoids redundancy and the **dummy variable trap** (perfect multicollinearity in regression).

■ Example

- Suppose we have a categorical variable ‘region’ with 4 categories: {north, east, south, west}.
 - Define $k - 1 = 3$ flag variables:
 - **north_flag**:
 - If region = north $\rightarrow 1$, Otherwise $\rightarrow 0$
 - **east_flag**:
 - If region = east $\rightarrow 1$, Otherwise $\rightarrow 0$
 - **south_flag**:
 - If region = south $\rightarrow 1$, Otherwise $\rightarrow 0$
 - The **west region** is the **reference category**.
 - If all three flags = 0, then region = west.

FLAG VARIABLES

■ Advantages

- Converts categorical data into numeric form, suitable for algorithms.
- Easy to interpret (1 = present, 0 = absent).

■ Limitations

- Can increase dataset dimensionality if categories are many.
- Requires careful choice of the **reference category**.

■ Applications in Data Mining

- **Regression models:** Encode categorical predictors as numeric.
- **Classification:** Enable algorithms to handle non-numeric inputs.
- **Clustering:** Represent categorical features in binary form.

TRANSFORMING CATEGORICAL VARIABLES INTO NUMERICAL VARIABLES

- Many data mining and machine learning algorithms require predictors to be numeric.
- However, real-world data often contains **categorical variables**
- To make such data usable, we must **transform categorical variables into numerical form**.
- Two common approaches:
 - Assign numeric codes directly (□ often wrong).
 - Use **flag variables (dummy/indicator variables)**

TRANSFORMING CATEGORICAL VARIABLES INTO NUMERICAL VARIABLES

■ The Wrong Approach: Arbitrary Numeric Coding

- Let's have a categorical variable **region** = {North, East, South, West}.
- If we assign numbers:
 - North = 1, East = 2, South = 3, West = 4
 - This creates a **false ordering**: West > South > East > North.
 - West is “three times closer” to South compared to North.
 - Algorithms wrongly assume **ordinal relationships** that do not exist.
 - This can lead to **biased or misleading model results**.

TRANSFORMING CATEGORICAL VARIABLES INTO NUMERICAL VARIABLES

- **The Correct Approach: Flag (Dummy) Variables**

- Instead of coding categories with numeric variable, we use **flag variables**:
- For **k categories**, create **k – 1 binary flags** (1 if present, 0 if absent).
- This avoids false ordering and makes the data suitable for numeric algorithms.

- **Exception: Ordered (Ordinal) Categorical Variables**

- Sometimes, categories have a **natural order (Ordinal Variables)**.
- Example: Survey responses → {Always, Usually, Sometimes, Never}.
- In such cases, numeric assignment is meaningful:
 - Always = 4, Usually = 3, Sometimes = 2, Never = 1

BINNING NUMERICAL VARIABLES

- Binning (also called discretization) is used to convert numerical variables into categories (bins or bands).
- Some algorithms prefer categorical rather than continuous predictors
- In such any numerical predictors are partitioned into bins or bands.
- For example, numerical variables ‘income’ value are portioned into - *low, medium, and high*.
- There are the following four common methods for binning numerical variables:

BINNING NUMERICAL VARIABLES

■ Equal Width Binning

- The numerical range is divided into k intervals of equal width.
- Example:
 - Suppose, $X=\{1,1,1,1,1, 2,2,11,11,12,12,44\}$, $k=3$.
 - Range = $44 - 1 = 43$, Width = $43 / 3 \approx 14.3$
 - Bins:
 - Low: $[1, 15.3)$ → contains most values $\{1,1,1,1, 1, 2,2,11,11,12,12\}$
 - Medium: $[15.3, 29.6)$ → contains no values
 - High: $[29.6, 44]$ → contains $\{44\}$
- Drawback: Outliers can distort bins, leaving some bins empty and others overloaded.

BINNING NUMERICAL VARIABLES

■ Equal Frequency Binning

- Each bin contains approximately the **same number of records**.
- If n = total records, and k = number of bins, then each bin has n/k records.
- **Example** (same dataset $X=\{1,1,1,1,1, 2,2,11,11,12,12,44\}$,
- $n = 12$, $k = 3 \rightarrow$ each bin = 4 records):
 - Low: {1,1,1,1}
 - Medium: {1,2,2,11}
 - High: {11,12,12,44}
- **Issue:** Identical values may fall into **different bins**, violating the principle that equal values should belong to the same category.

BINNING NUMERICAL VARIABLES

■ Binning by Clustering

- A clustering algorithm (e.g., **k-means**) is applied to group values.
- Automatically determines “natural” partitions in data.
- **Example:** Clustering the dataset $X=\{1,1,1,1,1, 2,2,11,11,12,12,44\}$,

Cluster 1: {1,1,1,1,1,2,2} (Low)

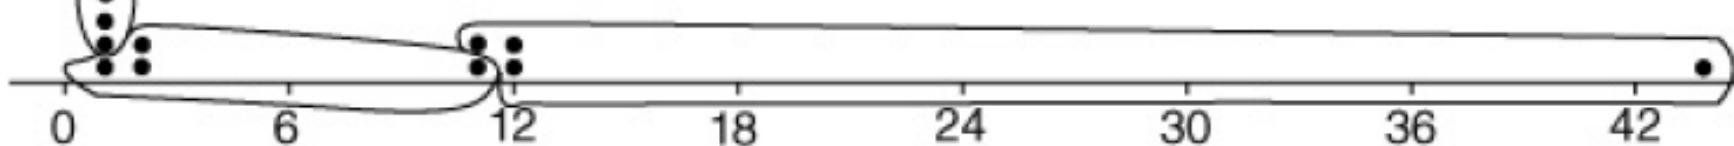
- Cluster 2: {11,11,12,12} (Medium)
- Cluster 3: {44} (High)
- Better than binning because it respects data distribution.

BINNING NUMERICAL VARIABLES

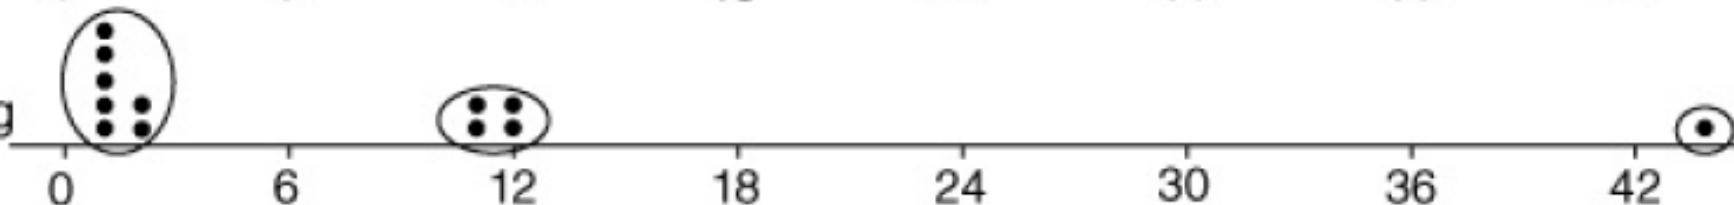
(a) Equal width



(b) Equal frequency



(c) k -Means clustering



BINNING NUMERICAL VARIABLES

■ Binning Based on Predictive Value

- Partitions are created with respect to the target variable (class label or outcome).
- Example: If 'income' value is used to predict "Loan Default," bins are chosen based on how values separate defaulters vs. non-defaulters.
- Best method when prediction is the goal.
- Considers target variable.

BINNING NUMERICAL VARIABLES

<u>Method</u>	<u>Approach</u>	<u>Pros</u>	<u>Cons</u>
Equal Width Binning	Divides range into equal widths	Simple	Distorted by outliers; empty bins possible
Equal Frequency Binning	Equal number of records per bin	Balances counts	Identical values may split
Clustering-based Binning	Uses clustering (e.g., k-means)	Respects data distribution	More complex
Predictive Value Binning	Based on target relationship	Improves prediction	Needs labeled data

RECLASSIFYING CATEGORICAL VARIABLES

- **Categorical variables** often contain **too many categories (levels or field values)**.
- Example: A variable like **State** in the U.S. has **50 categories**.
- Having too many categories can:
 - Increase model complexity and Reduce model accuracy.
 - Slow down computation.
 - Cause algorithms (e.g., logistic regression, decision trees) to perform **suboptimally**.
- To solve this, we use **Reclassification** (the categorical equivalent of binning numerical variables).

RECLASSIFYING CATEGORICAL VARIABLES

- **Reclassifying categorical variables** means combining or regrouping categories into **more meaningful groups** that are easier to analyze.
- Purpose: Reduce the number of categories while preserving or enhancing predictive power.
- Analogy: Just as binning numerical variables groups numbers into ranges, reclassifying categorical variables groups levels into categories.
- **Example**
 - **Original variable:** State (50 field values)
 - **Reclassified variable:** Region (5 categories)
 - Northeast, Southeast, North Central, Southwest, West
 - Instead of 50 unique values, the model now sees only 5.

RECLASSIFYING CATEGORICAL VARIABLES

- Methods of Reclassification

- Domain Knowledge-Based Grouping

- Use expert knowledge or business rules.
 - Example: Grouping states into **geographic regions** or **market zones**.

- Data-Driven Grouping

- Use clustering or similarity measures to group categories.
 - Example: Group products based on **sales patterns** or **customer demographics**.

- Target-Oriented Grouping

- Group based on relationship with the target variable.
 - Example: If predicting **loan default**, classify occupations into **High Risk**, **Medium Risk**, and **Low Risk**.

MOVING VARIABLES THAT ARE NOT USEFUL

- Not all collected variables contribute meaningfully to the analysis.
- Some variables may be redundant, irrelevant, or constant across the dataset.
- These variables must be removed; As, keeping such variables can lead to:
 - Increased computational cost
 - Noise in the data
 - Misleading results
- Two such variables include
 - unary variables and
 - variables that are very nearly unary.

MOVING VARIABLES THAT ARE NOT USEFUL

- **Unary Variables:** A unary variable is a variable that takes on only one value across all observation.
 - Example: If a dataset contains students from an all-girls private school, the variable `sex` would have only one value, “female”.
 - This variable has no effect on classification, regression, or clustering algorithms.
- **Nearly Unary Variables:** A variable that is almost constant, with one value dominating nearly all records.
 - Example: In a field hockey league dataset: 99.95% of players - female and 0.05% of players -male
 - Algorithms may treat such variables as essentially unary.
 - For example, a classifier could be >99.9% confident that a new record is “female,” rendering the variable uninformative for predictive modeling.

VARIABLES THAT SHOULD PROBABLY NOT BE REMOVED

- Analysts often remove problematic variables during preprocessing,
- But some imperfect ones may still hold valuable insights.
- Two common examples are:
 - Variables with a high proportion of missing values
 - Variables that are strongly correlated
- Removing them blindly can:
 - Reduce predictive power
 - Oversimplify models
 - Discard valuable signals
- Key Principle: Data preprocessing should be thoughtful, not automatic. Preserve potential sources of insight whenever possible.

VARIABLES THAT SHOULD PROBABLY NOT BE REMOVED

■ Variables with Many Missing Values

- Common Practice: Remove variables with $\geq 90\%$ missing values.
- Problem: Missingness itself may carry useful information and hidden patterns.
- Example: *donation_dollars* in a self-reported survey.
 - People who donate more may report proudly.
 - People who donate little may skip the question.
 - Thus, the missingness has a pattern and is not random.

VARIABLES THAT SHOULD PROBABLY NOT BE REMOVED

- Strategies Instead of Removal:
 - Create a flag variable (e.g., *donation_flag* = reported/not reported).
 - Captures missingness as potential predictor.
 - Impute missing values if the observed data are representative.
 - Use regression-based or decision tree-based imputation methods.
 - Evaluate representativeness of non-missing cases before imputing.
- Key Point:
 - Missingness may indicate characteristics, making it potentially predictive.
 - Avoid dropping variables simply because they contain lots of missing data.

VARIABLES THAT SHOULD PROBABLY NOT BE REMOVED

■ Strongly Correlated Variables

- Common Practice: Remove one of the variables.
- Problem: Removing correlated variables may discard important information.
- Example: *Precipitation vs. Beach Attendance*
 - Negatively correlated (more rain → fewer visitors).
 - Dropping one variable may ignore useful relationships.
- Alternative Approach:
 - Use PCA or Use algorithms that can handle correlation (e.g., tree-based models).
 - Transforms correlated predictors into uncorrelated principal components.
 - Retains information while avoiding multicollinearity.